

Tenacious-Bench Week 11 Report

Status report for the benchmark scaffold, scoring pipeline, and Path B critic preparation.

This PDF summarizes the current benchmark composition, the inter-rater agreement state, three worked scoring examples, and a direct status read on what is working now versus what still has to happen in Days 4 to 7.

1. Bench Composition

Tenacious-Bench v0.1 contains 210 tasks. The integrated cross-tab below shows counts across all three required axes at once: failure dimension, partition, and source mode. The partition target is hit exactly at 105 train, 63 dev, and 42 held-out tasks. The source-mode mix also lands on target to within integer rounding, with only a half-task rounding difference in the multi-LLM and hand-authored buckets.

Legend: the four source modes are mutually exclusive. Each task appears exactly once in the cross-tab, so any dimension-partition-source combination can be answered from a single lookup without double counting.

Integrated dimension x partition x source-mode cross-tab

Failure Dimension	train Trace	train Prog	train MLLM	train Hand	train Total	dev Trace	dev Prog	dev MLLM	dev Hand	dev Total	held_out Trace	held_out Prog	held_out MLLM	held_out Hand	held_out Total	All Total
Tone Style Drift	9	6	0	0	15	0	3	6	0	9	0	0	2	4	6	30
Hiring Signal Overclaiming	9	6	0	0	15	0	3	6	0	9	0	0	2	4	6	30
Bench Overcommitment	9	6	0	0	15	0	3	6	0	9	0	0	2	4	6	30
Competitor Gap Overclaiming	9	6	0	0	15	0	3	6	0	9	0	0	2	4	6	30
Public Signal Reliability	9	6	0	0	15	0	3	6	0	9	0	0	1	5	6	30
Scheduling Handoff Correctness	9	6	0	0	15	0	3	6	0	9	0	0	1	5	6	30
Icp Misclassification	9	6	0	0	15	0	3	6	0	9	0	0	1	5	6	30
All Dimensions	63	42	0	0	105	0	21	42	0	63	0	0	11	31	42	210

Target-versus-actual margins

Axis	Target Share	Target Count	Actual Count	Actual Share	Comment
Train / Dev / Held-out	50 / 30 / 20	105 / 63 / 42	105 / 63 / 42	50.0 / 30.0 / 20.0	Exact target split.
Trace-derived	30%	63.0	63	30.0%	Exact integer target.
Programmatic	30%	63.0	63	30.0%	Exact integer target.
Multi-LLM synthesis	25%	52.5	53	25.2%	Rounded up by 0.5 task.
Hand-authored adversarial	15%	31.5	31	14.8%	Rounded down by 0.5 task.