

2. Inter-Rater Agreement Results

Protocol: sample 30 tasks across the seven failure dimensions, label them once, wait 24 hours, then relabel blind against the same rubric dimensions. The acceptance rule is 80% agreement or better for each major rubric dimension. The current repository does not yet contain completed post-delay human relabel results, so this section reports status honestly rather than overstating maturity.

Current read: the repository has a complete agreement protocol and a deterministic scoring proxy, but not yet a finished human agreement matrix. That means the benchmark is operational for machine scoring, while the human-calibration layer should still be treated as open work before claiming that every rubric dimension is mechanically reliable.

Rubric Dimension	Metric	Target	Current Status	Interpretation
Required claims	Percent agreement	>= 80%	Pending human relabel	No claim yet; deterministic proxy is not counted as agreement.
Forbidden claims	Percent agreement	>= 80%	Pending human relabel	Will likely be most mechanically stable once relabeled.
Expected action	Percent agreement	>= 80%	Pending human relabel	Likely disagreement hotspot because phrase matching is stricter than human paraphrase tolerance.
Dimension guardrail	Percent agreement	>= 80%	Pending human relabel	Highest semantic judgment risk; should drive rubric revision if it misses.

Revision rule if any dimension lands below 80%: update the relevant rubric text in schema.json, document which task types caused the disagreement, rerun the 30-task relabel subset, and only then freeze the final wording for the held-out benchmark.

Interpretation: this makes the soft spots legible now rather than hiding them. Required-claims and forbidden-claims are expected to be the most mechanically stable dimensions, while expected-action and dimension-guardrail judgments are the most likely places where the human relabel loop will force rubric tightening.