

3. Worked Examples With Rubric Application

The three worked examples below show the evaluator applying the same deterministic scoring path to one trace-derived task, one programmatic task, and one hand-authored adversarial task. All three are deliberate partial-credit cases rather than perfect examples, which makes the scoring path more informative than a set of trivial passes.

Evaluator boundary for this report: these examples exercise the deterministic scoring layer only. Judge dimensions remain reserved in the schema for a later phase, but no live judge prompt/result is active in the current report, so that boundary is stated explicitly instead of implied.

tb-0001: Trace-derived example

Failure dimension: Tone Style Drift

Partition: train

Prospect: Orrin Labs (CTO)

Channel: email

Input evidence snapshot: funding signal confidence high; job-post signal confidence high; layoff signal confidence medium; leadership-change confidence low; AI maturity score 2; bench excerpt says 7 available for 6 requested in 7 days.

Candidate output: Question on Orrin Labs's engineering plans: Tenacious has engineers ready to deploy if the gap is real. Worth a discovery call next week?

Ground-truth field	Value
Required claims	engineers ready to deploy, question
Forbidden claims	A-player, bench, offshore, just checking in
Expected action	ask for a discovery call
Allowed evidence IDs	sig-tone_style_drift, bench_summary_v2, trace_log_week10

Scoring check	Pass?	Why it landed that way
required_claims_present	Yes	Matched mechanically in candidate output.
forbidden_claims_absent	Yes	Matched mechanically in candidate output.
expected_action_present	No	Missed by deterministic evaluator.
allowed_evidence_only	Yes	Matched mechanically in candidate output.
style_terms_absent	Yes	Matched mechanically in candidate output.
Final score	0.80	Pass threshold is 0.80 in the current evaluator.

Mechanical scoring path: The evaluator checks required-claim coverage, forbidden-claim absence, exact expected-action presence, an evidence overreach guard, and a dimension-specific guardrail. For this example the score is 0.80 because Failed expected_action_present. The schema still lists judge dimensions, but the committed example output here is scored by the deterministic evaluator only; there is no hidden judge prompt or judge override in this walkthrough.

tb-0064: Programmatic example