

4. Honest Status Assessment And Forward Plan

This is a healthy benchmark scaffold, not a finished benchmark publication. The repo already supports reproducible task generation, deterministic scoring, contamination checks, preference-pair preparation, and a first held-out comparison scaffold. The open risks are exactly the ones that matter at this stage: finishing the human agreement loop, replacing scaffolded corrections with a real trained critic, and spending the held-out evaluation budget only once the data pipeline is frozen.

Current cost posture remains conservative. The committed cost log is still local and deterministic only, and the held-out critic comparison adds only a small modeled step from \$0.0212 to \$0.0241 per task. The plan reserves \$2.50 of the \$10 challenge envelope for the sealed held-out evaluation pass, leaving the remainder for one small-model training run and one rerun if needed. The right kill criterion for Day 5 is simple: if the first live ORPO or SimPO run does not show stable training progress within the first thirty minutes, stop the run and debug data quality before spending more compute.

Working now

What is working	Evidence
Benchmark partitioning and composition	210 tasks in repo, exact 105 / 63 / 42 split, and source-mode mix within integer rounding of the 30 / 30 / 25 / 15 target.
Deterministic evaluator	Example tasks score reproducibly and the held-out scaffold can be compared numerically.
Contamination protocol	Held-out sealing protocol and contamination check are committed in the repo and used before freeze.
Path B training-data preparation	Preference-pair file exists and is ready for an ORPO or SimPO run.
Early ablation signal	Held-out scaffold shows 0.619 to 0.8571 average score lift with $p = 0.031$.

Blocked or still weak

What is not working yet	Why it matters
Inter-rater agreement not completed	The report cannot honestly claim that every rubric dimension clears the 80% bar.
No live ORPO or SimPO critic run yet	Current Delta A is a scaffolded correction benchmark, not a trained small-model critic.
Multi-LLM synthesis is opt-in	OpenRouter-backed synthesis remains optional, so some synthetic coverage is still router-generated rather than multi-model-generated.
Held-out evaluation budget not yet spent	The repo has reserved the eval phase, but the final sealed run still needs to happen within the ten-dollar envelope.

Days 4 to 7 plan

Day	Planned work	Exit criterion
Day 4	Finish path-specific reading pass, tighten preference-pair labels, and complete the 30-task human relabel cycle.	All agreement dimensions either clear 80% or trigger rubric revision.
Day 5	Run a small-backbone ORPO or SimPO critic on the train split using the prepared preference pairs.	Training shows stable loss movement inside 30 minutes; otherwise stop and inspect data quality.
Day 6	Run Delta A and Delta B on the sealed held-out split, plus cost and latency comparison.	Held-out evaluation completes within the reserved low-cost budget and yields a trustworthy comparison.
Day 7	Package final dataset, model artifact if any, memo, blog post, and community write-up.	All public artifacts point back to the same evidence graph and sealed metrics.