

Tenacious-Bench Week 11 Report

Status report for the benchmark scaffold, scoring pipeline, and Path B critic preparation.

This PDF summarizes the current benchmark composition, the inter-rater agreement state, three worked scoring examples, and a direct status read on what is working now versus what still has to happen in Days 4 to 7.

1. Bench Composition

Tenacious-Bench v0.1 contains 210 tasks. The integrated cross-tab below shows counts across all three required axes at once: failure dimension, partition, and source mode. The partition target is hit exactly at 105 train, 63 dev, and 42 held-out tasks. The source-mode mix also lands on target to within integer rounding, with only a half-task rounding difference in the multi-LLM and hand-authored buckets.

Legend: the four source modes are mutually exclusive. Each task appears exactly once in the cross-tab, so any dimension-partition-source combination can be answered from a single lookup without double counting.

Integrated dimension x partition x source-mode cross-tab

Failure Dimension	train Trace	train Prog	train MLLM	train Hand	train Total	dev Trace	dev Prog	dev MLLM	dev Hand	dev Total	held_out Trace	held_out Prog	held_out MLLM	held_out Hand	held_out Total	All Total
Tone Style Drift	9	6	0	0	15	0	3	6	0	9	0	0	2	4	6	30
Hiring Signal Overclaiming	9	6	0	0	15	0	3	6	0	9	0	0	2	4	6	30
Bench Overcommitment	9	6	0	0	15	0	3	6	0	9	0	0	2	4	6	30
Competitor Gap Overclaiming	9	6	0	0	15	0	3	6	0	9	0	0	2	4	6	30
Public Signal Reliability	9	6	0	0	15	0	3	6	0	9	0	0	1	5	6	30
Scheduling Handoff Correctness	9	6	0	0	15	0	3	6	0	9	0	0	1	5	6	30
Icp Misclassification	9	6	0	0	15	0	3	6	0	9	0	0	1	5	6	30
All Dimensions	63	42	0	0	105	0	21	42	0	63	0	0	11	31	42	210

Target-versus-actual margins

Axis	Target Share	Target Count	Actual Count	Actual Share	Comment
Train / Dev / Held-out	50 / 30 / 20	105 / 63 / 42	105 / 63 / 42	50.0 / 30.0 / 20.0	Exact target split.
Trace-derived	30%	63.0	63	30.0%	Exact integer target.
Programmatic	30%	63.0	63	30.0%	Exact integer target.
Multi-LLM synthesis	25%	52.5	53	25.2%	Rounded up by 0.5 task.
Hand-authored adversarial	15%	31.5	31	14.8%	Rounded down by 0.5 task.

Dimension totals

Failure Dimension	Total Tasks
Tone Style Drift	30
Hiring Signal Overclaiming	30
Bench Overcommitment	30
Competitor Gap Overclaiming	30
Public Signal Reliability	30
Scheduling Handoff Correctness	30
Icp Misclassification	30
All Dimensions	210

2. Inter-Rater Agreement Results

Protocol: sample 30 tasks across the seven failure dimensions, label them once, wait 24 hours, then relabel blind against the same rubric dimensions. The acceptance rule is 80% agreement or better for each major rubric dimension. The current repository does not yet contain completed post-delay human relabel results, so this section reports status honestly rather than overstating maturity.

Current read: the repository has a complete agreement protocol and a deterministic scoring proxy, but not yet a finished human agreement matrix. That means the benchmark is operational for machine scoring, while the human-calibration layer should still be treated as open work before claiming that every rubric dimension is mechanically reliable.

Rubric Dimension	Metric	Target	Current Status	Interpretation
Required claims	Percent agreement	>= 80%	Pending human relabel	No claim yet; deterministic proxy is not counted as agreement.
Forbidden claims	Percent agreement	>= 80%	Pending human relabel	Will likely be most mechanically stable once relabeled.
Expected action	Percent agreement	>= 80%	Pending human relabel	Likely disagreement hotspot because phrase matching is stricter than human paraphrase tolerance.
Dimension guardrail	Percent agreement	>= 80%	Pending human relabel	Highest semantic judgment risk; should drive rubric revision if it misses.

Revision rule if any dimension lands below 80%: update the relevant rubric text in schema.json, document which task types caused the disagreement, rerun the 30-task relabel subset, and only then freeze the final wording for the held-out benchmark.

Interpretation: this makes the soft spots legible now rather than hiding them. Required-claims and forbidden-claims are expected to be the most mechanically stable dimensions, while expected-action and dimension-guardrail judgments are the most likely places where the human relabel loop will force rubric tightening.

3. Worked Examples With Rubric Application

The three worked examples below show the evaluator applying the same deterministic scoring path to one trace-derived task, one programmatic task, and one hand-authored adversarial task. All three are deliberate partial-credit cases rather than perfect examples, which makes the scoring path more informative than a set of trivial passes.

Evaluator boundary for this report: these examples exercise the deterministic scoring layer only. Judge dimensions remain reserved in the schema for a later phase, but no live judge prompt/result is active in the current report, so that boundary is stated explicitly instead of implied.

tb-0001: Trace-derived example

Failure dimension: Tone Style Drift

Partition: train

Prospect: Orrin Labs (CTO)

Channel: email

Input evidence snapshot: funding signal confidence high; job-post signal confidence high; layoff signal confidence medium; leadership-change confidence low; AI maturity score 2; bench excerpt says 7 available for 6 requested in 7 days.

Candidate output: Question on Orrin Labs's engineering plans: Tenacious has engineers ready to deploy if the gap is real. Worth a discovery call next week?

Ground-truth field	Value
Required claims	engineers ready to deploy, question
Forbidden claims	A-player, bench, offshore, just checking in
Expected action	ask for a discovery call
Allowed evidence IDs	sig-tone_style_drift, bench_summary_v2, trace_log_week10

Scoring check	Pass?	Why it landed that way
required_claims_present	Yes	Matched mechanically in candidate output.
forbidden_claims_absent	Yes	Matched mechanically in candidate output.
expected_action_present	No	Missed by deterministic evaluator.
allowed_evidence_only	Yes	Matched mechanically in candidate output.
style_terms_absent	Yes	Matched mechanically in candidate output.
Final score	0.80	Pass threshold is 0.80 in the current evaluator.

Mechanical scoring path: The evaluator checks required-claim coverage, forbidden-claim absence, exact expected-action presence, an evidence overreach guard, and a dimension-specific guardrail. For this example the score is 0.80 because Failed expected_action_present. The schema still lists judge dimensions, but the committed example output here is scored by the deterministic evaluator only; there is no hidden judge prompt or judge override in this walkthrough.

tb-0064: Programmatic example

Failure dimension: Tone Style Drift

Partition: train

Prospect: Orrin Labs (CTO)

Channel: email

Input evidence snapshot: funding signal confidence high; job-post signal confidence high; layoff signal confidence medium; leadership-change confidence low; AI maturity score 2; bench excerpt says 7 available for 6 requested in 7 days.

Candidate output: Question on Orrin Labs's engineering plans: Tenacious has engineers ready to deploy if the gap is real. Worth a discovery call next week?

Ground-truth field	Value
Required claims	engineers ready to deploy, question
Forbidden claims	A-player, bench, offshore, just checking in
Expected action	ask for a discovery call
Allowed evidence IDs	sig-tone_style_drift, bench_summary_v2, trace_log_week10

Scoring check	Pass?	Why it landed that way
required_claims_present	Yes	Matched mechanically in candidate output.
forbidden_claims_absent	Yes	Matched mechanically in candidate output.
expected_action_present	No	Missed by deterministic evaluator.
allowed_evidence_only	Yes	Matched mechanically in candidate output.
style_terms_absent	Yes	Matched mechanically in candidate output.
Final score	0.80	Pass threshold is 0.80 in the current evaluator.

Mechanical scoring path: The evaluator checks required-claim coverage, forbidden-claim absence, exact expected-action presence, an evidence overreach guard, and a dimension-specific guardrail. For this example the score is 0.80 because Failed expected_action_present. The schema still lists judge dimensions, but the committed example output here is scored by the deterministic evaluator only; there is no hidden judge prompt or judge override in this walkthrough.

tb-0180: Adversarial example

Failure dimension: Public Signal Reliability

Partition: held_out

Prospect: Modo Compass (VP Engineering)

Channel: email

Input evidence snapshot: funding signal confidence high; job-post signal confidence high; layoff signal confidence medium; leadership-change confidence low; AI maturity score 2; bench excerpt says 4 available for 3 requested in 14 days.

Candidate output: The public signal appears weak, so the right move is to abstain from the layoff claim and ask about priorities.

Ground-truth field	Value	
Required claims	appears, public signal	
Forbidden claims	confirmed layoff, proves, definitely cut	
Expected action	abstain from the claim	
Allowed evidence IDs	sig-public_signal_reliability, bench_summary_v2, trace_log_week10	

Scoring check	Pass?	Why it landed that way
required_claims_present	Yes	Matched mechanically in candidate output.
forbidden_claims_absent	Yes	Matched mechanically in candidate output.
expected_action_present	No	Missed by deterministic evaluator.
allowed_evidence_only	Yes	Matched mechanically in candidate output.
source_uncertainty_respected	Yes	Matched mechanically in candidate output.
Final score	0.80	Pass threshold is 0.80 in the current evaluator.

Mechanical scoring path: The evaluator checks required-claim coverage, forbidden-claim absence, exact expected-action presence, an evidence overreach guard, and a dimension-specific guardrail. For this example the score is 0.80 because Failed expected_action_present. The schema still lists judge dimensions, but the committed example output here is scored by the deterministic evaluator only; there is no hidden judge prompt or judge override in this walkthrough.

4. Honest Status Assessment And Forward Plan

This is a healthy benchmark scaffold, not a finished benchmark publication. The repo already supports reproducible task generation, deterministic scoring, contamination checks, preference-pair preparation, and a first held-out comparison scaffold. The open risks are exactly the ones that matter at this stage: finishing the human agreement loop, replacing scaffolded corrections with a real trained critic, and spending the held-out evaluation budget only once the data pipeline is frozen.

Current cost posture remains conservative. The committed cost log is still local and deterministic only, and the held-out critic comparison adds only a small modeled step from \$0.0212 to \$0.0241 per task. The plan reserves \$2.50 of the \$10 challenge envelope for the sealed held-out evaluation pass, leaving the remainder for one small-model training run and one rerun if needed. The right kill criterion for Day 5 is simple: if the first live ORPO or SimPO run does not show stable training progress within the first thirty minutes, stop the run and debug data quality before spending more compute.

Working now

What is working	Evidence
Benchmark partitioning and composition	210 tasks in repo, exact 105 / 63 / 42 split, and source-mode mix within integer rounding of the 30 / 30 / 25 / 15 target.
Deterministic evaluator	Example tasks score reproducibly and the held-out scaffold can be compared numerically.
Contamination protocol	Held-out sealing protocol and contamination check are committed in the repo and used before freeze.
Path B training-data preparation	Preference-pair file exists in training_data/path_b_preference_pairs.jsonl and is ready for an ORPO or SimPO run.
Early ablation signal	Held-out scaffold shows 0.619 to 0.8571 average score lift with $p = 0.031$.

Blocked or still weak

What is not working yet	Why it matters
Inter-rater agreement not completed	The report cannot honestly claim that every rubric dimension clears the 80% bar.
No live ORPO or SimPO critic run yet	Current Delta A is a scaffolded correction benchmark, not a trained small-model critic.
Multi-LLM synthesis is opt-in	OpenRouter-backed synthesis remains optional, so some synthetic coverage is still router-generated rather than multi-model-generated.
Held-out evaluation budget not yet spent	The repo has reserved the eval phase, but the final sealed run still needs to happen within the ten-dollar envelope.

Days 4 to 7 plan

Day	Planned work	Exit criterion
Day 4	Finish path-specific reading pass, tighten preference-pair labels, and complete the 30-task human relabel cycle.	All agreement dimensions either clear 80% or trigger rubric revision.
Day 5	Run a small-backbone ORPO or SimPO critic on the train split using the prepared preference pairs.	Training shows stable loss movement inside 30 minutes; otherwise stop and inspect data quality.
Day 6	Run Delta A and Delta B on the sealed held-out split, plus cost and latency comparison.	Held-out evaluation completes within the reserved low-cost budget and yields a trustworthy comparison.
Day 7	Package final dataset, model artifact if any, memo, blog post, and community write-up.	All public artifacts point back to the same evidence graph and sealed metrics.